

Culminating Activity

Aloba, Loyan Edrian | Dalmonte, Franklin | Delrio Martin Nicholas | Dorin, Homer Adriel | Rivera, Nyt

May 08, 2026

Abstract (@Loyan)

Provide a concise summary of the study including purpose, methods, results, and conclusions (150–250 words).

Chapter 1: Introduction

Background of the Study

The integration of Generative AI (GenAI) into software engineering has been aggressively marketed as an unprecedented paradigm shift in productivity. Industry narratives frequently promise that AI coding assistants will streamline workflows, accelerate output, and consequently elevate both financial compensation and developer well-being. However, emerging empirical research reveals a significant productivity trade-off inherent in GenAI adoption. Utilizing the SPACE framework for developer productivity, Afroz et al. (2025) found that while developers using AI tools experience localized spikes in activity and perceive themselves as operating faster, they do not necessarily produce better software or experience increased fulfillment. In fact, rigorous randomized controlled trials indicate that developers using early-2025 AI tools took 19% longer to complete complex tasks, despite subjectively believing they were working 20% faster—a cognitive disconnect establishing an industry-wide “productivity placebo” (Becker et al., 2025). This paradox is driven by severe cognitive and operational friction. As AI accelerates initial code generation, it shifts systemic bottlenecks downstream, increasing pull request review times by 91% and inflating the volume of bugs per developer (Faros AI, 2025). Developers are forced to spend more time debugging hallucinated code and navigating complex, AI-generated architectural debt, which has been shown to actively damage macroscopic software delivery performance and system stability (DORA Research Program, 2024). Consequently, this study investigates whether the financial rewards of AI adoption genuinely translate into higher job satisfaction, or if the persistent friction of modern coding creates a “GenAI Productivity-Satisfaction Tradeoff.” In this paradigm, developers may be caught in a cycle where hyper-compensation—driven by a reported 56% AI skill premium (PricewaterhouseCoopers, 2025)—acts as a retention mechanism despite high stress, a concept previously identified in occupational psychology as job embeddedness or “golden handcuffs” (Mitchell et al., 2001). As AI increasingly shifts the developer’s role from writing original code to reviewing and debugging AI output, the industry must question whether these tools actually improve the human experience of software engineering. Utilizing secondary data from the Stack Overflow 2024 Developer Survey (Stack Overflow, 2024), this study builds a predictive model to determine if AI tools genuinely improve job satisfaction, or if

developers are simply compensated more to endure increased friction.

Statement of the Problem

This study investigates the systemic impact of AI adoption on software developers. With Job Satisfaction as the primary target variable, the study asks: To what extent does the use of AI development tools predict a developer’s overall job satisfaction? How does workflow friction (time spent searching for answers and debugging) interact with AI usage to influence this satisfaction? Does a higher salary offset the negative impact of workflow friction when predicting overall job satisfaction?

Objectives of the Study

General Objective

To investigate the “GenAI Productivity-Satisfaction Tradeoff” by developing a predictive model that evaluates how AI tool adoption, financial compensation, and workflow friction interact to determine a software developer’s overall job satisfaction.

Specific Objectives

To determine the extent to which AI tool adoption serves as a predictor for both financial compensation and workflow friction (time spent searching and debugging).

To identify the primary predictors of developer job satisfaction using a Classification and Regression Tree (CART) machine learning model.

To evaluate the interaction between compensation and workflow friction, specifically testing if a higher salary offsets the negative predictive impact of AI-induced friction on overall job satisfaction.

Hypotheses (if applicable)

H0 (Null Hypothesis): AI usage, financial compensation, and workflow friction do not form a significant predictive model for developer job satisfaction, and no distinct tradeoff exists between these variables.

H1 (Alternative Hypothesis): The optimal predictive model for Job Satisfaction will demonstrate a significant tradeoff, where the positive predictive value of high financial compensation is canceled out by the negative predictive impact of AI-induced workflow friction.

Significance of the Study

This research benefits technology companies, engineering managers, and developers. By identifying the gap between productivity metrics such as salary and output, and human metrics such as satisfaction, organizations can re-evaluate how they integrate AI. It shifts the focus from simply buying AI licenses to actively reducing the systemic friction that causes developer burnout.

Scope and Limitations

The study analyzes secondary data from the Stack Overflow 2024 Developer Survey. The scope is strictly limited to respondents who identify as “Employed, full-time” to maintain a stable economic baseline.

A critical limitation is that survey data regarding compensation and satisfaction is highly susceptible to MNAR (Missing Not At Random) bias, as extreme earners and highly burned-out developers frequently skip these questions. To mitigate this, advanced algorithmic imputation using MICE via Classification and Regression Trees was utilized to predict and fill missing values rather than relying on listwise deletion.

This study does not attempt an exhaustive feature selection of the entire Stack Overflow dataset. Instead, the predictor variables (AI Usage, Salary, Time Spent Searching, Years of Experience, and Remote Status) were handpicked based on intuition and prior literature regarding developer productivity.

Chapter 2: Review of Related Literature (@Loyan & @Nythan)

Related Studies

Summarize previous research.

Related Literature

Recent work by Stephany and Teutloff (2024) demonstrates that the economic value of a skill is strongly determined by *skill complementarity*—the number, diversity, and value of other skills with which it can be combined. Their analysis of 962 skills reveals that AI skills are particularly valuable, increasing worker wages by 21% on average, primarily because they exhibit high levels of complementarity with diverse existing skill sets. This finding provides theoretical grounding for understanding why financial compensation emerges as a critical mediating variable in predicting developer satisfaction. Rather than being arbitrary or divorced from productivity metrics, the salary premium associated with AI adoption reflects structural economic value accruing through skill portfolio enhancement.

Conceptual Framework

At the core of this framework is the ultimate target variable: Job Satisfaction.

Rather than utilizing isolated linear regressions, the framework employs a Decision Tree architecture where predictor variables interact organically:

Primary Predictor: AI Tool Adoption is positioned as the central independent variable, hypothesized to simultaneously impact financial compensation and workflow friction.

Interacting Predictors: Financial Compensation (Salary) and Workflow Friction (Time Spent Searching) act as mediating variables within the model. Contextual Controls: Years of Experience and Remote Status serve as baseline variables to account for demographic and workplace variances.

The CART algorithm evaluates how these variables interact, specifically testing whether high compensation mathematically buffers the negative impact of AI-induced workflow friction when predicting final job satisfaction levels.

Chapter 3: Methodology

Research Design




Figure 1: Relationship between variables

This study uses a quantitative, descriptive-correlational design to examine how AI adoption, salary, and workflow friction (as measured by Time Spent Searching & Debugging in minutes) relate to developer job satisfaction.

Data Collection

The study uses secondary data from the Stack Overflow 2024 Developer Survey. The raw CSV is provided in the repository `data/survey_results_public.csv`; the analysis filters this file to full-time respondents (see `filter_columns.r`) and writes the filtered table to `outputs/data_filtered_columns.csv`. The survey schema is included as `data/survey_results_schema.csv` for a clearer understanding of the possible columns from the original survey questionnaire.

Compensation and satisfaction responses are vulnerable to MNAR (Missing Not At Random) bias, so the analysis uses MICE with Classification and Regression Trees instead of listwise deletion.

The final modeling set was limited to Salary, Uses_AI, Years_Exp, Job_Satisfaction, Is_Remote, and Time-Searching because these variables align directly with compensation, AI adoption, seniority, work arrangement, and workflow friction.

Data Preparation and Initial Variable Selection

The raw survey file was reduced with `filter_columns.r`, which filtered for full-time employees and built the analysis table. Salary was derived from `ConvertedCompYearly`, Years_Exp from `YearsCodePro`, Uses_AI from `AISelect`, Job_Satisfaction from `JobSat`, and Is_Remote from `RemoteWork`. The output was saved to `outputs/data_filtered_columns.csv`.

These columns were kept because they capture the financial, behavioral, and contextual factors needed for this analysis’ statement of the problem and objectives. Salary was prioritized as a key predictor because prior research on skill complementarity (Stephany & Teutloff, 2024) establishes that compensation reflects structural economic value rather than arbitrary scarcity premiums, making it a theoretically grounded mediating variable between AI adoption and job satisfaction.

Initial Data Checks and EDA

The dataset was checked with `check_data.r` to review missingness patterns, unique-value counts, and basic distributions in both the raw survey data and the filtered modeling table. Exploratory visualization was handled with `eda.rmd` to inspect the early relationships among AI adoption, salary, job satisfaction, experience, and time spent searching/debugging. This step served as a plausibility check before modeling.

Imputation and Modeling Workflow

The imputation step was handled in `main.r`, which creates or reuses `outputs/data_imputed_primary.csv`.

An initial exploratory run used `set.seed(42)` (a commonly used placeholder seed; see Appendix C). Because a single RNG draw can create fragile results, the pipeline includes an explicit seed-sensitivity check: imputations were re-run with seeds = `c(42, 100, 200, 300, 400)` using `cart_robustness_check.r`, and the resulting trees were compared by their cross-validated error (`xerror`). This procedure ensures we report a defensible baseline while keeping the imputation uncertainty transparent.

After the sensitivity check, subsequent script runs adopted `set.seed(200)` for the corrected baseline because it produced the lowest cross-validated error across the tested seeds (see Chapter 4 and `outputs/cart_mice_seed_averages.csv`).

Hyperparameter Shuffling

The CART stage then uses hyperparameter shuffling to compare `cp`, `minsplit`, and `maxdepth`. In this context, `cp` is the pruning strength and `xerror` is the cross-validated error reported by `rpart`. Higher `cp` values prune more aggressively and can remove all splits. The best-performing combination was `cp = 0.001`, `minsplit = 20`, and `maxdepth = 5` (see hyperparameter outputs), and under the corrected baseline `TimeSearching_Min` was the strongest predictor, followed by `Years_Exp` and `Salary`.

Robustness Checks

The robustness work was separated into two parts. First, the hyperparameter shuffling held the imputed data fixed while varying model settings, which isolates model-choice effects. Second, the seed-sensitivity check changed the imputation seed while holding the CART settings fixed, which isolates data-uncertainty effects.

The seed-sensitivity workflow is documented in `cart_robustness_check.r`, and its outputs are saved to `outputs/cart_mice_seed_runs.csv` and `outputs/cart_mice_seed_averages.csv`. The CP profile is saved to `outputs/cart_cp_profile.csv` and provides an additional check on pruning.

Chapter 4: Results and Discussion

Descriptive Statistics

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(readr)
```

```
df <- read_csv("outputs/data_imputed_primary.csv", show_col_types = FALSE)
```

```
summary_df <- df %>%
  summarise(
    n = n(),
    Job_Satisfaction_mean = mean(Job_Satisfaction, na.rm = TRUE),
    Job_Satisfaction_sd = sd(Job_Satisfaction, na.rm = TRUE),
    Salary_mean = mean(Salary, na.rm = TRUE),
    Salary_sd = sd(Salary, na.rm = TRUE),
    TimeSearching_Min_mean = mean(TimeSearching_Min, na.rm = TRUE),
    TimeSearching_Min_sd = sd(TimeSearching_Min, na.rm = TRUE),
    Years_Exp_mean = mean(Years_Exp, na.rm = TRUE),
    Years_Exp_sd = sd(Years_Exp, na.rm = TRUE),
    Uses_AI_rate = mean(as.numeric(as.character(Uses_AI)) == 1, na.rm = TRUE),
    Is_Remote_rate = mean(as.numeric(as.character(Is_Remote)) == 1, na.rm = TRUE)
  )
```

```
summary_df
```

```
## # A tibble: 1 x 11
##       n Job_Satisfaction_mean Job_Satisfaction_sd Salary_mean Salary_sd
##   <int>           <dbl>           <dbl>         <dbl>     <dbl>
## 1 39041             6.93             2.08      88450.  122471.
## # i 6 more variables: TimeSearching_Min_mean <dbl>, TimeSearching_Min_sd <dbl>,
## #   Years_Exp_mean <dbl>, Years_Exp_sd <dbl>, Uses_AI_rate <dbl>,
## #   Is_Remote_rate <dbl>
```

Visualization (@Martin)

```
# Potential visualization: a compact bar chart or table snapshot of summary_df.
```

The descriptive table above is computed from outputs/data_imputed_primary.csv.

The imputed analysis sample provides the descriptive base for the rest of the chapter. Salary, job satisfaction, time spent searching/debugging, and years of experience are the central quantities because they capture the compensation, workload, and career-stage dimensions of a software engineer in the industry.

Model Complexity and Pruning Stability — Complexity Profile (CP)

Before evaluating hyperparameter choices, we first establish the fundamental complexity–accuracy tradeoff through the CP profile. This analysis shows how the pruning threshold constrains tree size and prediction error.

```
cp_profile <- read_csv("outputs/cart_cp_profile.csv", show_col_types = FALSE)
```

```
cp_profile
```

```
## # A tibble: 6 x 5
##       CP nsplit `rel error` xerror   xstd
##   <dbl> <dbl>   <dbl> <dbl> <dbl>
## 1 0.0118     0     1     1.00 0.00876
## 2 0.00553    1    0.988 0.988 0.00862
## 3 0.00267    2    0.983 0.983 0.00861
## 4 0.00141    3    0.980 0.981 0.00855
## 5 0.00116    4    0.979 0.980 0.00854
## 6 0.001      5    0.977 0.979 0.00853
```

Potential visualization: line plot of xerror against CP from outputs/cart_cp_profile.csv.

The CP profile in `outputs/cart_cp_profile.csv` supports the same conclusion but from another angle. In a decision tree algorithm, the CP balances predictive accuracy against structural complexity, effectively preventing the model from making unnecessary, hyper-granular splits that lead to overfitting.

An analysis of the CP profile shows a clear stabilization point. The table in `outputs/cart_cp_profile.csv` reports `xerror` values of: 1.00002 at $CP = 0.0118428$ (`nsplit = 0`), 0.98821 at $CP = 0.0055319$ (`nsplit = 1`), 0.98347 at $CP = 0.0026676$ (`nsplit = 2`), 0.98070 at $CP = 0.00140918$ (`nsplit = 3`), 0.98004 at $CP = 0.00115968$ (`nsplit = 4`), and 0.97894 at $CP = 0.001$ (`nsplit = 5`). The largest reductions in error occur with the earliest splits; after $CP = 0.001$ the marginal improvements are very small.

As seen above the cross-validated error across different pruning thresholds reveals a clear stabilization point at $CP = 0.001$. As the tree begins to grow, the initial broad splits, which represent the macro-structural drivers of Experience and Salary, produce the most significant and reliable reductions in prediction error. However, as the complexity constraint is loosened to allow for deeper branches, the cross-validated error flattens out and changes only within a very narrow range.

This behavior makes the chosen tree complexity highly defensible: the model attains most of its predictive benefit with a small number of splits, capturing the principal signals (years of experience and salary) without overfitting to minor fluctuations in workflow friction (time spent searching/debugging) and/or the adoption of AI tools.

Hyperparameter Shuffling

With the CP stabilization point at 0.001 identified, we now examine which hyperparameter combination achieves the best balance of accuracy and interpretability.

```
shuffle <- read_csv("outputs/cart_hyperparam_shuffle.csv", show_col_types = FALSE)

best_shuffle <- shuffle %>%
  filter(status == "ok") %>%
  arrange(xerror, desc(nsplitted)) %>%
  slice(1)

best_shuffle
```

```
## # A tibble: 1 x 12
##   rank status    cp minsplitted maxdepth xerror nsplitted Uses_AI Salary
##   <dbl> <chr>  <dbl>    <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>
## 1     5 ok     0.001      20        5  0.981    125        0   215.
## # i 3 more variables: TimeSearching_Min <dbl>, Years_Exp <dbl>, Is_Remote <dbl>
```

Hyperparameter Shuffling results are taken from `outputs/cart_hyperparam_shuffle.csv`. The best usable setting is `cp = 0.001`, `minsplitted = 20`, `maxdepth = 5` (`xerror = 0.98105`, `nsplitted = 125`).

By contrast, larger `cp` values (e.g., 0.008–0.02) frequently produce `nsplitted = 0` — a textbook sign of over-pruning. Over-pruning happens when the pruning penalty (`cp`) requires a splitting gain larger than any candidate split provides, so `rpart` collapses the tree to its root. The practical effects are underfitting (the model cannot partition the data), loss of variable-importance information, and misleading simplicity: a pruned tree with `nsplitted = 0` may appear stable but offers no explanatory partitions. Use the CP profile and grid search to pick `cp` near the elbow (about 0.001), where `xerror` is low and `nsplitted > 0`, balancing interpretability and generalization.

Potential visualization: xerror by cp, with nsplitted shown as a point label or second axis.

Set Seed Sensitivity

Next, we verify that the main results are not artifacts of a single imputation draw by testing sensitivity across different random seeds.

```
seed_averages <- read_csv("outputs/cart_mice_seed_averages.csv", show_col_types = FALSE)

seed_averages %>% arrange(xerror)
```

```
## # A tibble: 5 x 8
##   seed nsplitted xerror Uses_AI Salary TimeSearching_Min Years_Exp Is_Remote
##   <dbl> <dbl>    <dbl>    <dbl> <dbl>          <dbl>    <dbl>    <dbl>
## 1   200    185  0.978     0     652.          2503.    1102.     1.23
## 2   100    192  0.979     0     504.          2445.    1111.     4.56
## 3    42    159  0.979   2.75   343.          2510.    1113.     2.18
## 4   300    166  0.979     0     502.          2487.     971.     1.76
## 5   400    186  0.979   0.632  684.          2503.     998.     3.14
```


Potential visualization: variable importance by seed, using outputs/cart_mice_seed_averages.csv.

The MICE averages in `outputs/cart_mice_seed_averages.csv` show that the model is not driven by a single imputation draw. Across the tested seeds, `TimeSearching_Min` stays the strongest signal, while `Years_Exp` and `Salary` remain secondary. `Uses_AI` stays weak or near zero in most runs, which suggests that the direct effect of AI adoption is not stable once missing data is re-generated.

In this study the lowest observed `xerror` across the tested seeds occurred at `set.seed(200)` (`xerror` about 0.9784), so the main results are reported using the imputed dataset generated with `set.seed(200)`. The full seed-run results are available in `outputs/cart_mice_seed_runs.csv` and summarized in `outputs/cart_mice_seed_averages.csv`.

Testing for Robustness

The corrected baseline does not support AI usage as a strong driver of job satisfaction. In the main tree, `TimeSearching_Min` is the dominant split variable, with `Years_Exp` and `Salary` following behind. This means the model is showing that workflow friction is the strongest signal in the fixed-imputation run.

The robustness checks explain why this matters. When the model settings are shuffled, the same pattern stays visible: the chosen tree still favors `TimeSearching_Min` under the best-performing settings, while AI usage remains weak and unstable. When the missing data is perturbed through the seed sensitivity check, the same broad conclusion holds. `TimeSearching_Min` stays dominant, `Years_Exp` and `Salary` remain secondary, and `Uses_AI` contributes little or nothing in most runs.

The real drivers of satisfaction appear to be workflow friction, years of experience, and salary. This is closer to Herzberg’s (1959) view that satisfaction rests on basic working conditions and compensation, not on a single new tool. AI may still matter in limited cases, but it is not the main driver of the tree. In the corrected baseline, the model points first to the time spent searching and debugging, then to experience and pay.

This means GenAI is not a universal fix for low morale. It behaves more like a conditional feature than a core predictor. The 56% wage premium for AI skills (PricewaterhouseCoopers, 2025) can still work like “golden handcuffs” (Mitchell et al., 2001), but the tree does not show AI usage itself as the central force shaping satisfaction.

Interpretation

Taken together, the evidence suggests that the main tree is stable in its structure at the chosen hyperparameters, but the importance ranking is sensitive to imputation. In the corrected baseline, `TimeSearching_Min` is the strongest predictor, with `Years_Exp` and `Salary` trailing behind. When the missing data are dealt with through different imputations, workflow friction remains the most influential variable, while `Uses_AI` stays weak and unstable. This shows a weak and non-robust direct link between AI use and job satisfaction, while workflow friction, years of experience, and salary remain the strongest and most consistent signals.

Chapter 5: Conclusion and Recommendations

Conclusion

This research investigated the systemic impact of AI adoption on software developers, specifically examining whether the psychological burden of AI-induced workflow friction negates the positive impacts of financial

compensation. Addressing the core objectives established at the outset of the study, the findings reveal that the direct link between AI tool adoption and overall job satisfaction is surprisingly weak. While the initial problem statement questioned how workflow friction interacts with AI usage, the data demonstrates that at a broad career level, developers look past this friction. This behavior is driven by a “productivity placebo,” where the perception of speed outweighs the actual workload (Becker et al., 2025). However, at a granular operational level, the daily time spent debugging remains a volatile source of frustration, confirming the warnings from Afroz et al. (2025) regarding the hidden effort of reviewing AI outputs. Regarding the initial inquiry of whether a higher salary offsets this friction, the findings point to a structural reality rather than a direct mathematical tradeoff. High compensation acts as “golden handcuffs” (Mitchell et al., 2001), successfully keeping developers in their roles. This wage premium reflects the structural economic value of AI skills through complementarity with diverse existing competencies (Stephany & Teutloff, 2024), making the financial incentive rational from both employer and employee perspectives. The daily annoyance of technical debt is endured because the financial compensation makes it rational to stay, rather than the friction being directly “canceled out” by the pay. These findings align directly with the observations by Afroz et al. (2025): faster task completion through GenAI does not inherently represent true progress. Organizations should avoid the assumption that integrating AI coding assistants will automatically yield higher productivity or a happier workforce.

The data clearly indicates that GenAI can create “spurious productivity,” where surface-level acceleration hides the deeper, manual effort required to review and validate AI outputs (Afroz et al., 2025).

Consequently, the alternative hypothesis (H1), which states that there is a significant tradeoff in which high financial compensation is canceled out by AI-induced friction, is definitively rejected (See Chapter 4 for Set Seed Sensitivity and Testing for Robustness subsections, where `Uses_AI` remains weak or near zero in most runs). The anticipated tradeoff was not reflected in the analysis of the dataset, because the predictive power of AI usage is not reflected in the different robustness checks. Overall job satisfaction is driven more by foundational factors like years of experience and competitive pay, while workflow friction operates independently as an expected daily challenge. AI adoption remains a highly conditional modifier, not a universal driver of a developer’s professional well-being.

Recommendations

The following recommendations are intended for future researchers who want to extend or replicate this study’s method. They are grounded in the patterns observed in Chapter 4, especially the sensitivity of the tree to imputation and the recurring importance of workflow friction over AI usage.

1. Compare imputation strategies directly. MICE should remain the primary benchmark because it handles mixed variable types well and preserves uncertainty across multiple draws, but future studies can test KNN-based imputation as a sensitivity check. If a clustering-based alternative is explored, it should be used cautiously and justified against the data structure, since not all clustering methods handle mixed categorical and numeric variables cleanly. The main goal is to determine whether the same findings hold under a different missing-data treatment.
2. Use the imputation comparison to test robustness, not just accuracy. A future paper could run the same modeling pipeline under MICE, KNN, and one additional baseline method, then compare the downstream effects on tree structure, xerror, and variable importance. That would show whether the conclusions are stable or whether they depend too heavily on one imputation choice.

3. Extend the analysis to salary as a second outcome. A useful follow-up would be to build a parallel model for Salary and then compare it with the Job Satisfaction model. Since the results and discussion suggest that compensation and satisfaction are related but not identical outcomes, this would let future researchers check whether the same predictors appear in both models and whether AI tool adoption has a similar or different role depending on the response variable.
4. Make AI adoption more explicit in future models. The results section shows that AI usage was not the dominant driver in the final tree, so future researchers should still keep it in the model but consider stronger operationalizations, such as frequency of use, task-specific use, or usage intensity. This would help determine whether the null result comes from the construct itself or from the way it was measured.
5. Report seed and model stability in every replication. Because imputation and tree fitting are both stochastic, future work should report the seed used for the main run, the full seed-sensitivity table, and the selected hyperparameters. A small appendix of alternate runs is often enough to show that the result is not a one-seed artifact.
6. If time allows, compare CART with at least one alternative model. CART was appropriate here because the goal was interpretability, but future studies could compare it with random forest or a regularized regression model to check whether the same predictors appear consistently across different model families.
7. Add a cross-outcome association check when the workflow is expanded. The results and discussion make it reasonable to expect overlap between the Salary and Job Satisfaction models, so future researchers can formally test that relationship using a concordance check, correlation screen, or two-model comparison. That would be a clean way to see whether AI adoption, experience, and workflow friction line up across outcomes without forcing the idea into the current study’s already-complete results section.

References

- Afroz, S., Feng, Z., Kimura, K., Trinkenreich, B., Steinmacher, I., & Sarma, A. (2025). Developer productivity with GenAI. arXiv. <https://arxiv.org/html/2510.24265v1>
- Becker, J., Rush, N., Cunningham, T., & Rein, D. (2025). Measuring the impact of early-2025 AI on experienced open-source developer productivity. arXiv. <https://arxiv.org/abs/2507.09089>
- DORA Research Program. (2024). 2024 Accelerate state of DevOps report. Google Cloud. <https://dora.dev/research/2024/dora-report/>
- Faros AI. (2025). The AI productivity paradox research report. <https://www.faros.ai/blog/ai-software-engineering>
- Herzberg, F. (1959). The motivation to work. John Wiley & Sons.
- Mitchell, T. R., Holtom, B. C., Lee, T. W., Sablinski, C. J., & Erez, M. (2001). Why people stay: Using job embeddedness to predict voluntary turnover. *Academy of Management Journal*, 44(6), 1102-1121. <https://doi.org/10.2307/3069391>
- PricewaterhouseCoopers. (2025). 2025 Global AI jobs barometer. <https://www.pwc.com/gx/en/issues/artificial-intelligence/job-barometer/2025/report.pdf>

Stack Overflow. (2024). 2024 Developer Survey. <https://survey.stackoverflow.co/2024/>

Stephany, F., & Teutloff, O. (2024). What is the price of a skill? The value of complementarity. *Research Policy*, 53(1), 104927. <https://doi.org/10.1016/j.respol.2023.104927>

Appendices

Appendix A: Code (@Nythan)

The analysis code is included in the code chunks throughout this R Markdown document.

Appendix B: Data Source Stack Overflow. (2024). 2024 Developer Survey. Retrieved from <https://survey.stackoverflow.co/2024/>

Appendix C: Notes on Random Seeds

Random seeds initialize pseudorandom number generators (PRNGs) so that stochastic processes (like MICE imputations or algorithmic subsampling) can be reproduced exactly. A seed is not a model parameter — it only fixes the RNG state for reproducibility. Because different seeds can produce slightly different imputations or model splits, reporting the seed and testing sensitivity across several seeds is good practice for transparency.

Why use 42? The number 42 is a cultural reference (Douglas Adams' *The Hitchhiker's Guide to the Galaxy*) and has become a de-facto, tongue-in-cheek default in many code examples. There is no statistical or computational reason to prefer 42 over any other integer — it is purely conventional.

Recommendation: report the seed used for the published baseline, include the full seed-sensitivity results (as found in `outputs/cart_mice_seed_averages.csv`), and treat the seed choice as part of the reproducibility metadata rather than a hyperparameter that needs theoretical justification.